


Observation and analysis of the multicore performance impact on scientific applications



Tyler A. Simon^{1,*} and James McGalliard²

¹*Computer Sciences Corporation, Greenbelt, MD, U.S.A.*

²*FEDSIM, Greenbelt, MD, U.S.A.*

SUMMARY

With the proliferation of large multicore high-performance computing systems, application performance is often negatively affected. This paper provides benchmark results for a representative workload from the Department of Defense High-performance Computing Modernization Program. The tests were run on a Cray XT-3 and XT-4, which use dual- and quad-core AMD Opteron microprocessors. We use a combination of synthetic kernel and application benchmarks to examine the cache performance, MPI task placement strategies and compiler optimizations. Our benchmarks show performance behavior similar to that reported in other studies and sites. Dual- and quad-core tests show a run-time performance penalty compared with single-core runs on the same systems. We attribute this performance degradation to a combination of L1 to main memory contention and task placement within the application. Copyright © 2009 John Wiley & Sons, Ltd.

Received 6 April 2009; Accepted 23 June 2009

KEY WORDS: multicore; benchmark; high-performance computing; memory contention; MPI

1. INTRODUCTION

Generally speaking, application runtimes on large high-performance computing (HPC) systems with dual- or quad-core chips are slower when compared, core for core, with single-core chip performance. This paper provides a multilevel approach to analyzing application performance on multicore architectures. This paper presents an analysis of the results of a variety of currently used applications by examining latencies within the processor, compiler optimizations and the MPI communication layer. We contrast the results with similar studies and synthetic test results. Furthermore, the authors suggest alternatives for addressing the performance loss that the ‘multicore era’ has introduced to the scientific computing community.

*Correspondence to: Tyler A. Simon, Computer Sciences Corporation, Greenbelt, MD, U.S.A.

†E-mail: tyler.simon@nasa.gov



The synthetic kernel and application benchmark tests were run at the U.S. Army Engineer Research and Development Center, Department of Defense (DoD) Supercomputer Resource Center (ERDC DSRC) at the U.S. Army Corps of Engineers in Vicksburg, Mississippi. The ERDC DSRC is part of the DoD high-performance computing modernization program (HPCMP). These tests show performance behavior and challenges similar to those reported in other studies and at other sites, for example that contention for memory degrades the performance of multicore processors—in this case Cray cluster systems using dual- and quad-core AMD processors. Low-level memory read/write time measurements using synthetic kernels are consistent with wall-clock runtimes measured using application software representative of HPCMP science and engineering workloads.

This paper is organized as follows:

- Section 2 describes the Sapphire and Jade benchmark systems that the authors used.
- Section 3 reviews a sample of prior studies of multicore processor performance.
- Section 4 presents our analysis of Sapphire and Jade test results and some contrasting results from other sites.
- Section 5 discusses some of the practical implications of our results.
- Section 6 presents our conclusions.

2. BENCHMARK SYSTEMS

For this study, there were two systems that were upgraded with multicore AMD processors, a Cray XT3 that went from single-core to dual-core and a Cray XT-4 that went from dual-core to quad-core, both of which are located at the U.S. Army ERDC in Vicksburg and as of November 2008 are ranked internationally at 75th and 45th, respectively [1]. An overview of each system is provided below.

2.1. Sapphire

The Sapphire system includes the following and is rated at 42.6 Tflops peak performance.

- 4096 2.6 GHz 64-bit Dual-Core AMD-Opteron (8192 computation cores).
- 2 GB memory per core.
- Total Disk Storage 374 TB, Fiber Channel Raid.
- 3-D torus using a HyperTransport link to a dedicated Cray SeaStar.
- Custom MPI message-passing implementation.

2.2. Jade

The Jade system includes the following and is rated at 36.2 Tflops peak performance.

- 2152 2.1 GHz 64-bit Quad-Core AMD-Opteron (8608 computation cores).
- 2 GB memory per core.
- Total Disk Storage 379 TB, Fiber Channel Raid.
- 3-D torus using a HyperTransport link to a dedicated Cray SeaStar2.




Figure 1. Opteron design [2].

2.3. System design

Both the Sapphire Cray XT-3 and Jade Cray XT-4 use multicore Opteron processors. Figure 1 is a high-level picture of the Opteron design, whose features for both systems include the following:

- Multiple cores—two for the XT-3 and four for the XT-4.
- Independent 64 kB level 1 (L1) instruction, 64 kB L1 data and 1 MB L2 caches.
On-chip Translation Lookaside Buffer (not shown).
- Shared DDR access paths to main memory off the chip.
- Catamount proprietary operating system on the XT3 and UNICOS/lc 2.0.x on the XT4.

The benchmark results reported later in the paper demonstrate some of the impacts of these design features.

It is worth noting the following with regard to the AMD Opteron processors. Both the L1 and L2 caches are private to each core, and access to main memory, which may be local to the node or on a remote node in the cluster, is shared by all cores. Further, with the Opteron the L2 cache is treated as an eviction cache or ‘victim’ cache, which means it receives least recently used (LRU) data from L1 and is probed by L1 when there is a local miss. When the data cannot be found in either L1 or L2, it is copied from the main memory into L1. Thus managing main memory to L1 cache bandwidth is important on the Opteron, especially when considering that each core must access main memory independently, that is without the help of a shared cache.

Below, we also present benchmark results that compare Sapphire to two other DoD systems, Midnight (at the Arctic Region Supercomputing Center) and JVN (at the Army Research



Laboratory). The key characteristics of these systems are as follows:

Midnight

- Sun Fire X4600.
- 716 2.6 GHz 64-bit Dual-Core AMD-Opteron (1432 cores).
- 4 GB memory per core.
- Four cores per node.
- 64k instruction and 64k data Level 1 cache per core, not shared.
- 1 MB Level 2 cache per core, not shared.
- 4X Voltaire Infiniband commercial network mesh.

JVN

- Linux Networkx Evolocivity II.
- 2048 3.6 GHz 64-bit Single-core Intel EM64T Xeon (2048 cores).
- 2 GB memory per core
- Two cores per node.
- 12k instruction and 16k data Level 1 cache per core, not shared.
- 2 MB Level 2 cache per core, not shared.
- Myrinet network mesh.

The benchmark results and these specifications yield insights into the multicore performance and strategies for alleviating the performance bottlenecks.

3. BACKGROUND

The problem of multicore memory contention has received considerable attention in the literature. Dongarra *et al.* [3] present a concise assessment of the multicore effect on large-scale HPC and put forth a clear agenda for addressing this ‘multicore revolution’ from a software development perspective. In general, our results are consistent with others’. Following are some points from prior studies.

Alam *et al.* [4] pointed out that on the XT3 at the Oak Ridge National Laboratory, the contention for the memory link shared by multiple cores limits performance and that communications within a core vs. across multiple cores showed an 8–12% benefit. Additionally, they found that appropriate MPI task and memory placement can yield a 25% performance benefit.

Regarding the MPI performance on multicore systems, Chai *et al.* [5] found that around 50% of the messages are intra-node vice inter-node, so performance gains can be seen by focusing on intra-node optimization. In addition, Chai suggests that multicore aware applications and message-passing algorithms optimized for data locality can improve performance by up to 70%.

Finally, Levesque *et al.* report in [6] that for the NAS benchmarks, the average dual-core degradation was 27%. We compare our NAS kernel results with Levesque’s specific results in Figure 5.

Additional noteworthy information from his paper is that the general atomic and molecular electronic structure system (GAMESS) code showed a 2 to 3% degradation when running in dual-core mode vs. single-core mode—one of the application codes that we also tested. Further results show that the dual-core performance degradation is directly proportional to degradation in memory



bandwidth and that dual core results have observed memory bandwidth almost exactly one-half of single-core results for the same tests.

4. PERFORMANCE ANALYSIS

Prior to our discussion of application runtimes, we will examine the difference between memory latencies as we multiply cores on nodes.

4.1. Opteron memory performance

In order to measure the memory performance, we composed a synthetic kernel that executes a series of memory reads and writes while summing the elapsed time and iterating over progressively increasing memory stride sizes based on the model developed in Hennessy and Patterson [7]. The results in Figure 2 show measured performance plateaus corresponding to the size of the Opteron caches.

In Figure 2 the horizontal axis is the memory stride size as a power of 2 and the vertical axis is the memory latency—memory read and write time in nanoseconds.

The benchmark reads and writes will occasionally miss in the cache, and memory bandwidth will impact the performance when it does so. The results show that running the benchmark on two cores causes an approximate doubling of memory R + W time for strides between 2^8 and 2^{13} bytes, and less significant but consistent degradation for larger stride sizes [8].

The red line represents the performance of a single copy of the program running on one core of the dual-core Opteron, the green line represents performance when both cores on the same chip

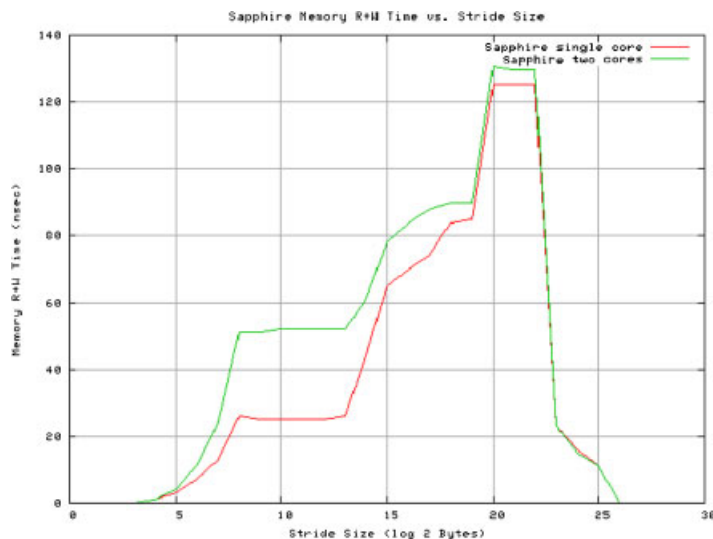


Figure 2. HPCMP XT3 memory stride benchmark results.

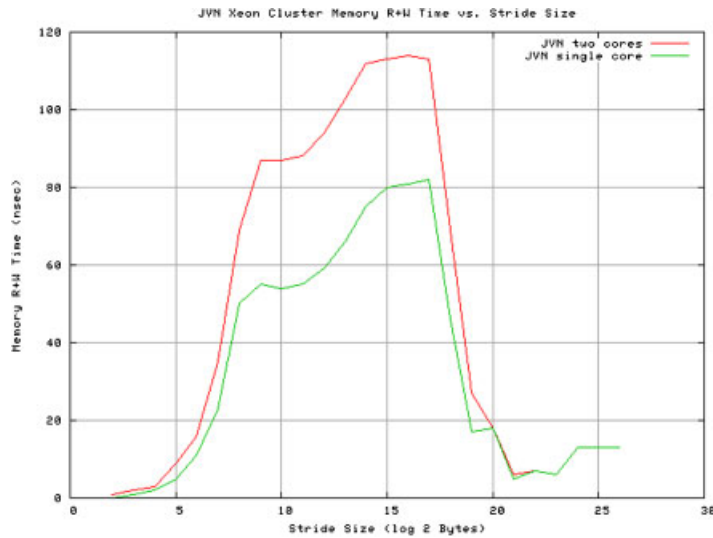


Figure 3. Intel Xeon memory stride benchmark results.

are running individual copies of the same memory stride kernel. As stride size increases from 2^{15} (64 k, the same size as the L1 cache) to 2^{20} (1M, the same size as the L2 cache), memory latency reflects gradual filling of the L2 cache, followed by a steep jump (latency degradation) when the L2 cache is nearly full.

There is significant performance degradation, on the order of doubling latency between 2^8 and 2^{14} byte strides, comparing the dual- to the single-core data. The low latency when the code is running in the L1 cache and significant degradation when running dual-core over the same memory stride show that it is worth focusing on the performance of the L1 cache. Note that the AMD cache hierarchy differs from dual-core Intel chips because its L2 cache is a victim cache. The LRU data is evicted from L1 into L2, then L1 copies from main memory if the data is not found in L2. This underscores the fact that L1 to main memory bandwidth is worth managing for large multicore Opteron installations.

There is a performance plateau for both single- and dual-core tests between about 2^{20} and 2^{22} (1 MB and 4 MB stride). This reflects the performance of access to main storage before the prefetch mechanism is engaged.

Latency drops steeply above a 2^{22} stride. The XT3 possesses both hardware and software prefetch facilities, and the results indicate that the prefetch mechanism is engaged and dominates the performance for the longest memory strides in the test.

We compared the XT3 using Opteron processors to the JVN cluster that uses the Intel Xeon chip, which has two cores on two sockets that share main memory on the node, as shown in Figure 3.

In Figure 3 the axes are the same as in Figure 2, except that the peak latency is slightly lower (120 vs. 140 ns). As with the Opteron results in Figure 2, two cores have increased latency in comparison to single-core when running the synthetic memory stride code. The degraded Xeon two core performance is most pronounced at the stride lengths that correspond to the L1 cache size.

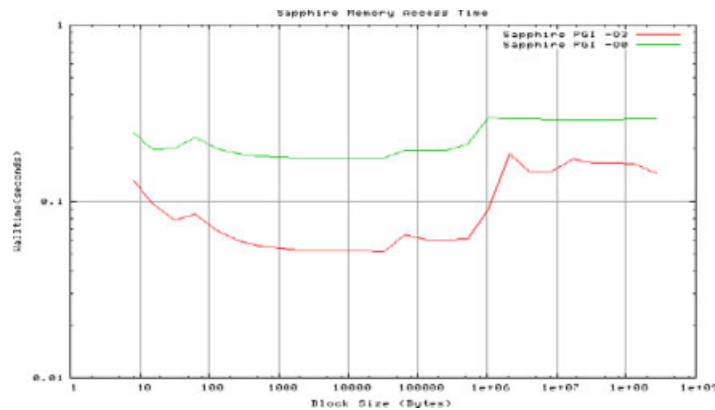


Figure 4. Sapphire XT3 memory access time.

The two-core/dual-core memory penalty persists up to about the same stride length, 2^{18} , on both systems. The Cray Sapphire, with two cores per processor or socket, experiences contention at the socket level; the Linux Networkx JVN, with one core per processor or socket, does not experience contention at that level; both experience contention at the board or node level. Both experience contention for main storage.

The performance of both systems is similar up to about a 256 byte stride.

Up to about a 2^{18} stride, Sapphire has faster memory latency than JVN. The custom SeaStar network mesh on Sapphire is faster than the commercial Infiniband mesh on JVN.

Like the AMD Opteron, the Xeon processors have hardware and software prefetch facilities, and the prefetch mechanism is engaged and dominates the performance for the longest memory strides in the test. The prefetch mechanism engages just after the program overruns the L2 cache, showing precipitous performance improvement above a 2^{18} stride. This mechanism is engaged at smaller strides on JVN compared with Sapphire (2^{18} vs. 2^{22} or 128 k vs. 4 MB).

This kernel does not test inter-node contention.

4.2. Compiler optimization and memory performance

We were curious about the impact of compiler optimizations on multicore performance. In addition to the standard compiler optimization levels, the `-O3` optimization automatically blocks and unrolls loops and prefetches data into the cache. Figure 4 shows these results for the PGI compiler on the Sapphire system. The memory kernel test is the same as in Figures 2 and 3, but block size rather than stride length varies.

In Figure 4 the horizontal and vertical axes are logarithmic scales. The vertical axis is memory latency measured in seconds; the horizontal axis is the size of the blocks written to and read from memory. The green line in the figure shows memory access performance with optimization suppressed (`-O0`). The red line shows memory access with the aggressive compiler optimization, `-O3`. `-O3` results are significantly faster than `-O0`, by a factor on average of about 3:1. This is due to managing the L1 cache.



Notice that there are performance shifts near the 65k and 1 MB block sizes that correspond to the L1 and L2 cache sizes. Once each level of cache fills up and the job starts to access the next level of cache or main storage, the performance degrades. Note also that the performance improves between the smallest block sizes and about 50 000 bytes. This suggests that as use of the cache increases to a certain point, the latency improves. More data is reused from the cache, so the average latency drops.

4.3. Effect of main memory bandwidth

We ran the NAS kernel benchmarks on the XT3 and XT4—the same tests, Opteron processors and Cray clusters as Levesque *et al.* ran (Section 3, [6]). Figure 5 compares Levesque's results with ours, which were quite close except for BT.

Levesque's paper proposed a simple model for predicting the performance impact of multicore performance based on the insight that virtually all degradation is due to contention for memory bandwidth. '[P]redicted dual-core performance across a wide range of memory access patterns is correctly accounted for using the simplest possible model for memory performance—namely, that the effective latency of the memory subsystem and bandwidth+latency for the L2 cache remain the same regardless of how many cores are used, but the main memory bandwidth is cut in half.' Our results, using the same NAS kernels shown in Figure 5 but decomposing the performance into L1/L2, L1/MM, L2/MM and wall clock runtime, support this insight—all elements are strongly correlated with main memory bandwidth. Figure 6 shows this.

In his paper, Levesque argues that the Translation Lookaside Buffer and cache are *not* the source of the dual-core memory latency penalty. We did not have counter data, as did he. For Opteron, where the TLB and caches are *not* shared, it is plausible that what Levesque claims is true. For the single-core Xeons (JVN), by the same token, TLB and cache are also not shared because they are on the processor chip. It is not clear that this is an advantage. Later Intel processors, such as Harpertown, do share the L2 cache across cores—the tradeoff is that there can be increased locality

Code	Time Increase for Dual Core (Levesque)		Time Increase for Dual Core (Jade)	Delta Jade vs. Levesque Large
	Large Pages	Small Pages		
BT	109%	112%	142%	1.266
SP	130%	139%	159%	1.141
LU	138%	138%	131%	0.952
CG	110%	110%	106%	0.966
MG	146%	145%	139%	0.956

Figure 5. Comparison of Levesque and Jade NAS Kernel benchmark results—single vs. dual-core performance degradation.




Figure 6. Ratios of single-core vs. dual-core performance metrics.

and higher L2 cache hit rates when there is not strictly enforced processor affinity vs. the increased contention for the shared cache.

4.4. MPI intra- and inter-node communication performance

Prior analysis indicates that two key drivers of multicore performance are memory bandwidth and message passing— ‘We observed from the AMD architectural discussion that when excluding messaging performance, the primary source of contention when moving from single-core to dual-core is memory bandwidth’ [5]. We also tested message-passing performance, as this section discusses. Our results are consistent with Levesque’s.

We are interested in what effect multicore chips have on the MPI performance. MPI is an open-system standard and various MPI implementations are available, so portability across platforms is simplified. (Contrast MPI with Cray *shmem*, a proprietary interface that allows for logically shared distributed memory access. Although the hardware specificity of Cray *shmem* can improve the performance compared with the more general MPI interface, the cost can be a more difficult migration to the next generation platform.)

4.4.1. MPI performance on Sapphire and Midnight systems

The next series of graphics compare the performance of the Sapphire system and Midnight, a Opteron-based Sun Linux cluster at the Arctic Region Supercomputing Center (another DoD HPCMP site). The test shows an average bandwidth from sending a 500 MB buffer between ranks in a point-to-point fashion using `mpi_send` and `mpi_recv`.

In all of the following graphs that depict results in three dimensions, the apparent vertical axis represents memory throughput measured in Megabytes per second. There are two apparent horizontal axes, pitched toward the upper left and upper right. These axes represent the MPI rank of the source and destination of an MPI message. (System administrators can rank or label their nodes for



use by MPI in any arbitrary fashion. The important point to grasp is that in cells where the source and destination ranks are the same [along the diagonal of the apparent horizontal plane], the MPI message is passed between cores on the same node. As pointed out by Chai *et al.* [5], intra-node communication is much faster than inter-node communication.)

Throughput is indicated in the graphs both by color (yellow is highest, blue is lowest) and by apparent topographic peaks and valleys. The range of throughputs in each graph is different, but can be interpreted from the legends on either side of the graph. The range of ranks in each graph is also different, but can be interpreted from the labels along the two apparent horizontal axes. In each case, the range of ranks is the number of cores in that test.

Figure 7 shows MPI throughput on HPCMP's Sapphire system using eight cores on four nodes (dual-core processors). Overall, the performance is fairly even and high. There is less than a 10% difference between the minimum and maximum throughput. There is some fall-off in throughput where there is off-node communication, for example Rank 0 communicating with Rank 7. Throughput is somewhat higher where MPI communication is on the node, that is along the diagonal—Rank 0 communicating with Rank 0, Rank 1 communicating with Rank 1 and so on.

Figure 8 results are similar to Figure 7, except that the number of cores and nodes has doubled. This 16 core test has even more consistent message-passing performance than the 8 core test. This test also ran on Sapphire for 32, 64 and 128 cores, and the results were consistent—high bandwidth across all pairs of MPI ranks, that is both intra-node and inter-node performances. We attribute the consistent performance of the Sapphire system to Cray's design—the Cray SeaStar network mesh and its MPI implementation [9].

For contrast, Figures 9–12 show results from Midnight, a Linux cluster at the Arctic Region Supercomputing Center. Using more generic, less highly optimized components (compared with Cray) but the same Opteron chips yielded significantly lower and more variable performances. The format of the Midnight graphs is the same as for Sapphire.

Midnight is a Linux Networkx cluster system consisting of dual-core 2.6 GHz Opteron processors and a total of 1432 cores, as explained above. As the number of cores and nodes increases (Figures 9–12), the overall performance degrades quickly [7].

Midnight cores	Midnight nodes	Min MPI throughput	Max MPI throughput
16	4	1000 MB s ⁻¹	1400
32	8	500	1000
64	16	500	850
128	32	480	680

There is an immediate, consistent and rapid degradation of MPI throughput as more cores are added to the test. One twenty-eight-core MPI throughput performance is one half of 16 core performance. There is a pronounced ridge along the diagonal—on node and off node bandwidth can differ by $\frac{1}{3}$. One can get 30% better performance by putting the MPI rank on the same node. The difference between intra-node communication and inter-node communication increases as we increase the number of cores. This is common with large multicore-cluster installations.

Midnight is a traditional Linux cluster—not as optimized as the Cray Sapphire. Recalling the Sapphire performance, there was much higher performance and much more consistent performance

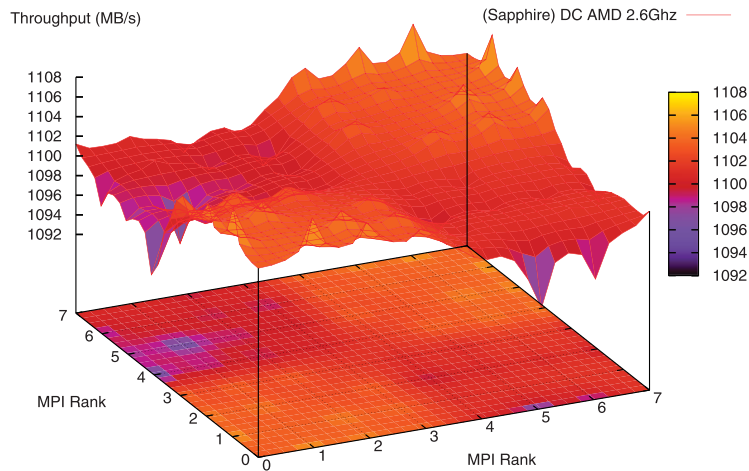


Figure 7. Sapphire MPI message throughput by MPI rank, 8 cores on 4 nodes.

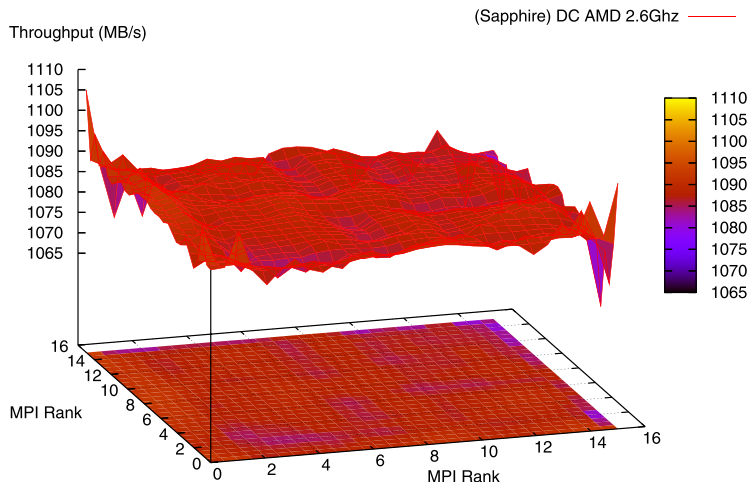


Figure 8. Sapphire MPI message throughput by MPI rank, 16 cores on 8 nodes.

across the range in number of cores in the test and across pairs of MPI ranks, compared with the much more variable and slower performance of Midnight.

4.4.2. MPI rank reordering

In view of the impact of MPI rank ordering shown in the preceding tests, we also ran tests to change the MPI rank assignments. On the XT3 the default Round Robin assignment of MPI processes to ranks has the effect of spreading processes across multiple nodes, which can dilute spatial reference

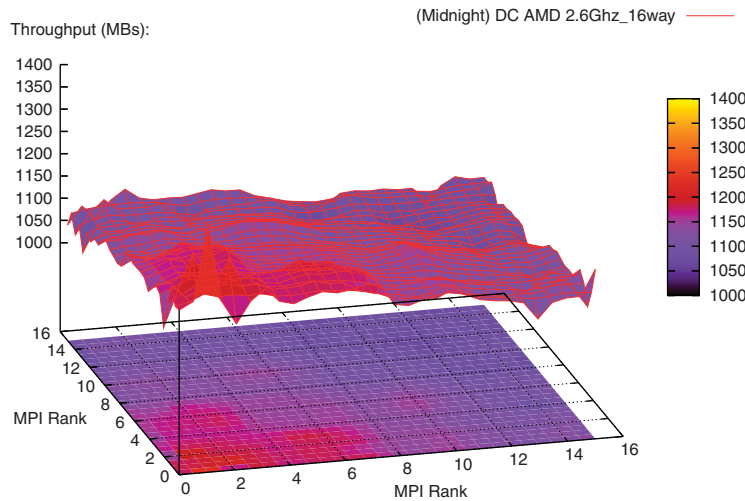


Figure 9. Midnight MPI message throughput by MPI rank, 16 cores on 4 nodes.

locality and degrade performance. The symmetric multiprocessing (SMP) MPI option increases reference locality, increases the proportion of intra-core vice inter-core message traffic and so improves performance on certain codes. Figure 13 illustrates the results of running the GAMESS application with both placement policies on Sapphire; we can see that as the core count increased the performance, gains are reduced. This is likely caused by inter-node communication overhead.

The following section shows the benchmark results for complex application software packages. The results of the memory stride kernel (Section 4.1), NAS Kernel (Section 4.3) and MPI (Section 4.4) low-level benchmarks predict that memory and message passing will cause dual-core performance degradation for such complex applications, and our results bear out this prediction.

4.5. HPCMP TI application benchmark results

Each year the DoD assesses its computational workload and develops a benchmark suite which it uses for acquisitions. The HPCMP releases this package to vendors for benchmarking purposes [10]. Figure 14 provides a short description and references for the codes included in the study.

This section shows Sapphire (XT3) and Jade (XT4) results using this HPCMP TI benchmark suite. For Figures 15–18, the horizontal axis reflects the number of cores running in the test. Each test was run in single- and dual-core modes for Figures 15 and 16. In dual-core mode, both cores on a single chip ran the same benchmark code.

In Figures 17 and 18, each test was run in dual- and quad-core mode. In quad-core mode, all four cores on the same processor ran the benchmark code. The vertical axis is the ratio between the dual-core and single-core runtimes (Figures 15 and 16) or quad-core and dual-core runtimes (Figures 17 and 18).

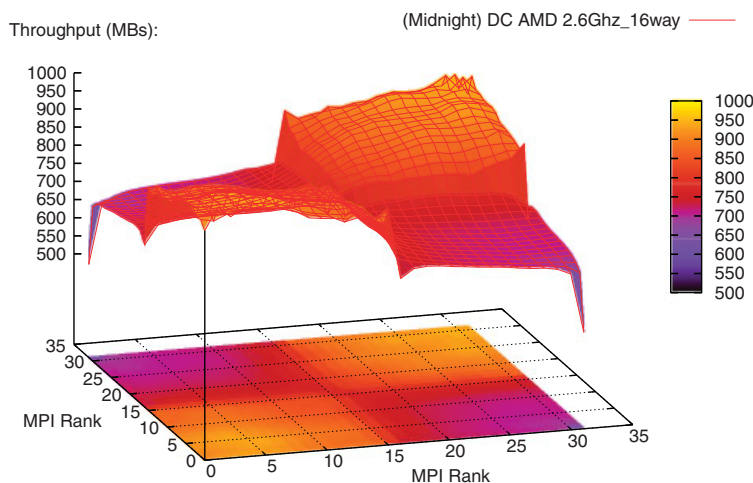


Figure 10. Midnight MPI message throughput by MPI rank, 32 cores on 8 nodes.

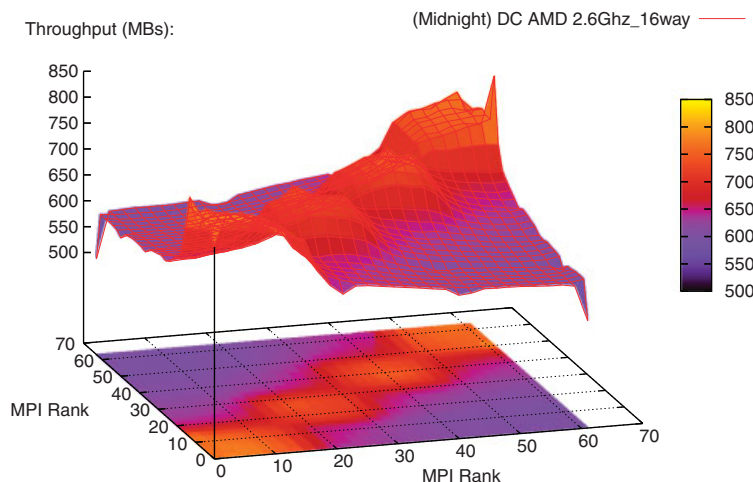


Figure 11. Midnight MPI message throughput by MPI rank, 64 cores on 16 nodes.

Figure 15 shows results for 32 to 256 cores and Figure 16 shows 256 to 768 cores. For all codes except GAMESS, dual-core mode runtimes were slower than single-core mode. The GAMESS outlier results were based on a test using the Cray shm interface between cores, all the other tests used MPI. Cray shm takes advantage of specific low-level hardware features to achieve better performance. MPI is an open standard that would lose its generality and portability if it used low-level hardware calls. In addition, GAMESS has two types of tasks—a data server process and a compute process. These two processes can work in concert on the same node and give better a performance than the other codes.

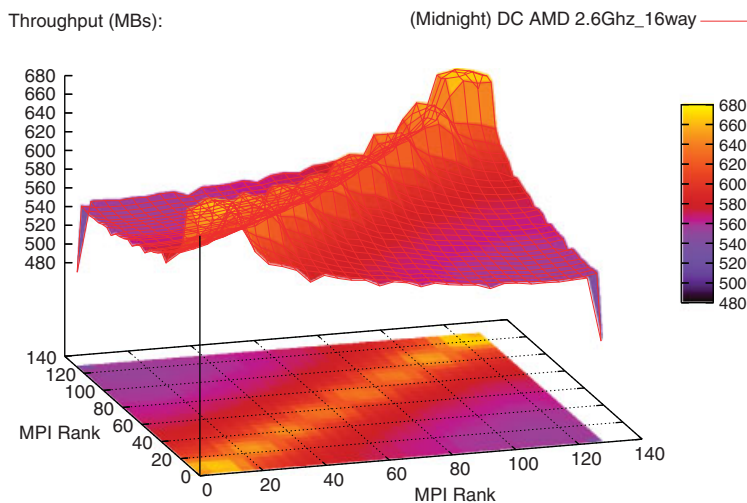


Figure 12. Midnight MPI message throughput by MPI rank, 128 cores on 32 nodes.

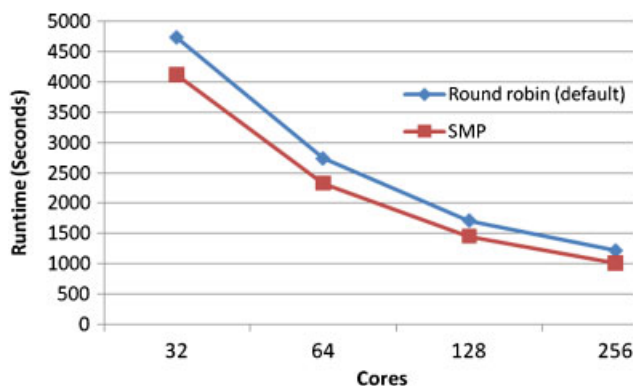


Figure 13. GAMESS XT3 MPI rank reorder method results.

Figures 17 and 18 show the results on the XT4 for dual- and quad-core runs.

The consistent performance degradation in multiple single-core vs. dual-core comparisons is not displayed in our dual- vs. quad-core results for these applications. The most important reason is that the dual- and quad-core tests were run 1 year apart, with upgrades to the XT4 hardware and system software in between the two tests, with significant performance improvements attributable to these upgrades.

In addition, the specific characteristics of the application program impact the performance and can dominate multicore performance degradation. A closer examination of these codes should reveal the specific reasons for their multicore performance patterns.



Code	Acronym	Comments & References
Adaptive Mesh Refinement	AMR	* Solves 3D Navier-Stokes Equation
Air Vehicles Unstructured Solver	AVUS	* From Air Force Research Laboratory * Finite Volume Unstructured Grid Euler/Navier-Stokes Solver * www.erdhpc.mil/hardSoft/Software/avus
Shock Physics	CTH	* From Sandia National Laboratories * Large deformation, strong shock-wave, solid mechanics code * www.erdhpc.mil/hardSoft/Software/cth
General Atomic and Molecular Electronic Structure System	GAMESS	* Quantum Chemistry, Restricted Hartree-Fock * Communication and Memory Intensive * www.msg.ameslab.gov/games/gamess.html * www.erdhpc.mil/hardSoft/Software/games * Uses dual processes well suited for dual-core processors
Hybrid Coordinate Ocean Model	HYCOM	* Ocean General Circulation Model * oceanmodeling.rsmas.miami.edu/hycom/
Improved Concurrent Electromagnetic Particle-in-Cell	ICEPIC	* From Air Force Research Laboratory * Simulates collisionless plasma physics on a Cartesian grid
Out of Core Matrix Solver	OOCORE	* Dense Linear Solver * Memory and I/O Intensive
Overflow2	O'FLOW	* Computational Fluid Dynamics
Weather Research and Forecast Model	WRF	* An Advanced Mesoscale Forecast and Data Assimilation System * Scalable across a range of problems and many computer platforms * www.docstoc.com/docs/4903066/THE-WEATHER-RESEARCH-AND-FORECAST-MODEL-(WRF)

Figure 14. HPCMP benchmark application codes.

(Incidentally, it is not unusual to see the runtime performance increase as the number of cores increases—a fixed-size problem that can run in parallel on multiple cores will complete faster. Increasing core counts for a fixed-size problem can also improve the cache hit rates. As the problem is divided among more and more processors, the working set shrinks and will run more inside the cache(s) and faster. Parallel processing and working sets that fit in the cache can both improve performance for fixed-size problems.)

5. DISCUSSION

These results have practical implications now and over the middle term.

Multicore and processor-aware code optimization. It is not sufficient to simply run an application on more cores to achieve better performance. Users should examine their codes and consider restructuring them to increase locality, increase intra-node communications, assign MPI ranks to promote spatial locality, use compiler optimizations and make other multicore aware changes.

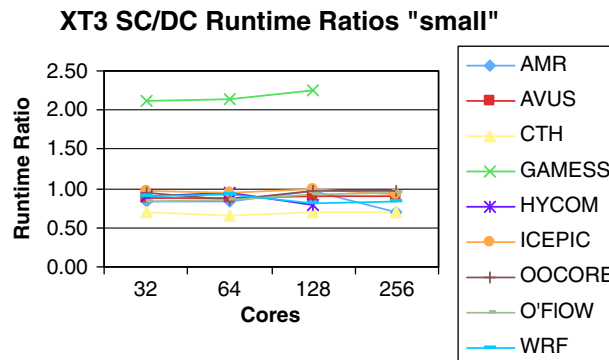


Figure 15. Single- vs. dual-core ratios, 'Small' XT3 cores.




Figure 16. Single- vs. dual-core ratios, 'Larger' XT3 cores.

Insight into multicore specifically and processor and cluster design generally can help application performance. Microprocessor designs such as Intel's and AMD's have different application performance implications; cluster designs such as Cray's and Sun's also have diverse performance impacts.

Where runtimes are the critical constraint and processor utilization is less critical, reducing the number of active cores per chip can help to improve the performance. Some schedulers can support single-core runs.

Tuning MPI placement, such as by rank reordering, can help application performance.

Users can improve multicore performance by optimizing their codes to increase reference locality, e.g. loops and data structures that fit within known cache sizes, and by using compiler options. Compilers can affect multicore performance differently than single-core performance. Revisit your compiler optimizations or look to upgrade to compilers designed for multicore environments.

Parallel API alternatives. Think about using the OpenMP application programming interface because it helps to manage resources locally—on the shared memory processors and nodes. Local resource management was not necessary when every task was equivalent but it is now.

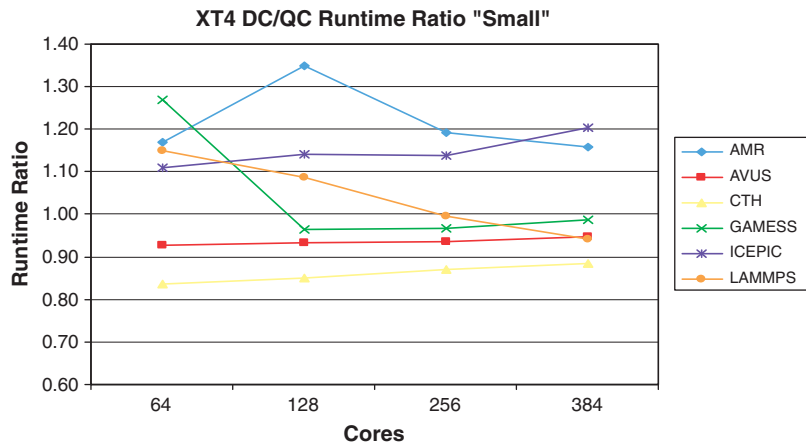


Figure 17. Dual- vs. quad-core ratios, 'Small' XT4 cores.

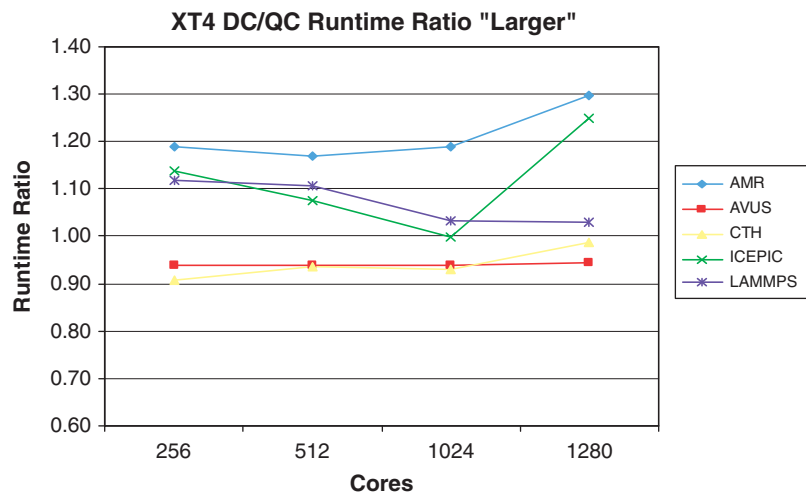


Figure 18. Dual- vs. quad-core ratios, 'Larger' XT4 cores.

MPI placement can impact performance, but does not automatically take advantage of any local resources, as OpenMP can. MPI does not manage tasks, but Open MP can manage user-level tasks for you.

Market/technology futures. Quad-core+ multi-socket nodes will likely exacerbate the bandwidth contention issue both to main memory and on-chip over the middle term. Many applications will experience significant negative multicore processor performance impacts unless optimized to account for them.



The processor manufacturers have announced plans to use new technologies to mitigate the multicore performance problems. Hybrid disruptive heterogeneous multiprocessor environments may help the scientific computing community. Perhaps a combination of a traditional Linux cluster combined with cell processors will provide a more cost-effective solution than systems built solely on commodity microprocessors. However, we do not know yet if these approaches will work.

6. CONCLUSIONS

The memory stride kernel shows degraded dual-core runtimes compared with single-core. Single- vs. dual-core impacts on cache and main memory communications are dominated by memory bandwidth problems. Our results suggest that memory bandwidth contention is the single most important cause of multicore performance degradation.

Using alternate compiler optimizations with the memory stride kernel suggests that more aggressive optimizations such as -O3 can improve the performance and do so by improving the use of the L1 cache.

The NAS benchmark results show that multicore performance impacts many applications, but the impact varies according to the applications' characteristics. The results also show that the performance impact is most strongly correlated with the main memory bandwidth, consistent with Levesque's results.

The MPI kernel results show that intra-node performance is consistently faster than inter-node performance. Our results indicate that a highly optimized, custom design such as the Cray SeaStar can support better performance than unmodified commercial products such as the Voltaire Infini-band. Finally, MPI placement is under the application programmer's control and can be used to increase locality and performance.

Taken together, these results support others' published conclusions that multicore processors can significantly harm the performance, but that users can mitigate that impact by careful design and optimization of their codes, and that the processor manufacturing community recognizes and is responding to these multicore problems in their announced plans for future products.

ACKNOWLEDGEMENTS

Mr Simon appreciates the discussion of the above material and the comments provided by Bill Ward, Sam Cable, Paul Bennett, Tom Oppe (ERDC), Roy Campbell (DoD HPCMP), Brian Waldecker (AMD), Harvey Wasserman (NERSC) and Alan Minga (Cray).

REFERENCES

1. Available at: <http://www.top500.org/> [6 April 2009].
2. Available at: <http://www.amd.com> [6 April 2009].
3. Dongarra J, Gannon D, Fox G, Kennedy K. The impact of multicore on computational science software. *CTWatch Quarterly* 2007; 3(1):1–6.
4. Alam SR, Barrett RF, Kuehn JA, Roth PC, Vetter JS. Characterization of scientific workloads on systems with multicore processors. *IEEE International Symposium on Workload Characterization*, San Jose, CA, October 2006; 225–236.
5. Chai L, Gao Q, Panda DK. Understanding the impact of multicore architecture in cluster computing: A case study with Intel dual-core system. *International Symposium on Cluster Computing and the Grid*, Rio de Janeiro, Brazil, 2007.



6. Levesque J, Larkin J, Foster M, Glenski J, Geissler G, Whalen S, Waldecker B, Carter J, Skinner D, He H, Wasserman H, Shalf J, Shan H, Strohmaier E. Understanding and mitigating multicore performance issues on the AMD opteron architecture. *Paper LBNL-62500*, Lawrence Berkeley National Laboratory, 7 March 2007.
7. Hennessy J, Patterson D. *Computer Architecture: A Quantitative Approach* (3rd edn). Morgan Kaufmann: San Mateo, CA, 2003; 516.
8. Chandra D, Guo F, Kim S, Solihin Y. Predicting inter-thread cache contention on a chip multi-processor architecture. *Proceedings of International Symposium on High Performance Computer Architecture (HCPA)*, San Francisco, 12–13 February 2003.
9. Brightwell R, Pedretti K, Underwood KD. Initial performance evaluation of the Cray SeaStar Interconnect. *Proceedings of the 13th Symposium on High Performance Interconnects (HOTI'05)*. IEEE: New York, 2005; 51–57.
10. Available at: <http://www.hpcmo.hpc.mil/bizopps/TI/> [6 April 2009].